

# Garbage Classification — Data Preprocessing & Inspection

In this notebook, we will perform the full **data preprocessing and inspection** steps to make the dataset ready for modeling.

We aim to satisfy the following criteria:

- Load & inspect the dataset ( $\approx 13.9k$  images)
- Confirm the 6 classes ( plastic , metal , glass , cardboard , paper , trash )
- Verify class balance ( $\approx 2,300$ – $2,500$  images per class)
- Detect (and optionally remove or flag) duplicates
- Confirm image sizes and color channels (expected:  $256 \times 256$ , 3 channels RGB)
- Ensure labels align correctly with image files
- Split into train / validation / (test) sets, with stratification
- Normalize / standardize pixel values (record method)
- (Optional) Set up data augmentation
- Build a pipeline or loader to ensure a batch can go through a baseline CNN

We'll break this into sections.

---

## Step 0: Download Kaggle Garbage Images Dataset

```
In [ ]: from google.colab import files
# Upload your kaggle.json
# Only needs to be done once
# If you have not uploaded kaggle.json file here before,
# follow instructions below on acquiring kaggle.json
files.upload()
!mkdir -p ~/.kaggle
!mv kaggle.json ~/.kaggle/
!chmod 600 ~/.kaggle/kaggle.json
```

Choose Files No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving kaggle.json to kaggle.json

```
In [ ]: !pip install -q kaggle

import os

# Make sure the Kaggle API key is available
if not os.path.exists("/root/.kaggle/kaggle.json"):
```

```

print("UPLOAD KAGGLE.JSON FILE ABOVE")
print("WATCH VIDEO I (AGUSTIN) SENT IN GROUP CHAT ON GETTING KAGGLE.JSON")

# Create data directory if not exists
os.makedirs("data", exist_ok=True)

# Download + unzip only if not already present
if not os.path.exists("data/Garbage_Dataset_Classification"):
    !kaggle datasets download -d zlatan599/garbage-dataset-classification -p data/
    !unzip -q data/garbage-dataset-classification.zip -d data/
    print("Dataset downloaded and extracted!")
else:
    print("Dataset already exists, skipping download.")

```

Dataset URL: <https://www.kaggle.com/datasets/zlatan599/garbage-dataset-classification>

License(s): MIT

Downloading garbage-dataset-classification.zip to data

0% 0.00/121M [00:00<?, ?B/s]

100% 121M/121M [00:00<00:00, 1.64GB/s]

Dataset downloaded and extracted!

## Step 1: Setup & Imports (install if not already done)

In [ ]: `%pip install imagededup`

Collecting imagededup

Downloading imagededup-0.3.3.post2-cp312-cp312-manylinux\_2\_24\_x86\_64.manylinux\_2\_28\_x86\_64.whl.metadata (8.0 kB)

Requirement already satisfied: torch in /usr/local/lib/python3.12/dist-packages (from imagededup) (2.8.0+cu126)

Requirement already satisfied: torchvision in /usr/local/lib/python3.12/dist-packages (from imagededup) (0.23.0+cu126)

Requirement already satisfied: Pillow>=9.0 in /usr/local/lib/python3.12/dist-packages (from imagededup) (11.3.0)

Requirement already satisfied: tqdm in /usr/local/lib/python3.12/dist-packages (from imagededup) (4.67.1)

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.12/dist-packages (from imagededup) (1.6.1)

Requirement already satisfied: PyWavelets in /usr/local/lib/python3.12/dist-packages (from imagededup) (1.9.0)

Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-packages (from imagededup) (3.10.0)

Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (1.3.3)

Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (0.12.1)

Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (4.60.1)

Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (1.4.9)

Requirement already satisfied: numpy>=1.23 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (2.0.2)

Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (25.0)

Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (3.2.5)

Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.12/dist-packages (from matplotlib->imagededup) (2.9.0.post0)

Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn->imagededup) (1.16.2)

Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn->imagededup) (1.5.2)

Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.12/dist-packages (from scikit-learn->imagededup) (3.6.0)

Requirement already satisfied: filelock in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (3.20.0)

Requirement already satisfied: typing-extensions>=4.10.0 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (4.15.0)

Requirement already satisfied: setuptools in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (75.2.0)

Requirement already satisfied: sympy>=1.13.3 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (1.13.3)

Requirement already satisfied: networkx in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (3.5)

Requirement already satisfied: jinja2 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (3.1.6)

Requirement already satisfied: fsspec in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (2025.3.0)

Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (12.6.77)

Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in /usr/local/lib/p

```

python3.12/dist-packages (from torch->imagededup) (12.6.77)
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (12.6.80)
Requirement already satisfied: nvidia-cudnn-cu12==9.10.2.21 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (9.10.2.21)
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (12.6.4.1)
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (11.3.0.4)
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (10.3.7.77)
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (11.7.1.2)
Requirement already satisfied: nvidia-cusparse-cu12==12.5.4.2 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (12.5.4.2)
Requirement already satisfied: nvidia-cusparselt-cu12==0.7.1 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (0.7.1)
Requirement already satisfied: nvidia-nccl-cu12==2.27.3 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (2.27.3)
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (12.6.77)
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (12.6.85)
Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (1.11.1.6)
Requirement already satisfied: triton==3.4.0 in /usr/local/lib/python3.12/dist-packages (from torch->imagededup) (3.4.0)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-packages (from python-dateutil>=2.7->matplotlib->imagededup) (1.17.0)
Requirement already satisfied: mpmath<1.4,>=1.1.0 in /usr/local/lib/python3.12/dist-packages (from sympy>=1.13.3->torch->imagededup) (1.3.0)
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.12/dist-packages (from jinja2->torch->imagededup) (3.0.3)
Downloading imagededup-0.3.3.post2-cp312-cp312-manylinux_2_24_x86_64.manylinux_2_28_x86_64.whl (318 kB)

```

318.6/318.6 kB 9.6 MB/s eta 0:00:00

Installing collected packages: imagededup

Successfully installed imagededup-0.3.3.post2

```

In [ ]: # Required Libraries
import os
from pathlib import Path
from collections import Counter
import random
import hashlib

from PIL import Image
import numpy as np
import matplotlib.pyplot as plt

# For splitting
from sklearn.model_selection import train_test_split

# For dedup
from imagededup.methods import PHash

```

## Step 2: Define dataset root & discover classes

```
In [ ]: dataset_root = Path("data/Garbage_Dataset_Classification/images")

assert dataset_root.exists(), f"Dataset root not found: {dataset_root}"

# List subdirectories as candidate classes
classes = [d.name for d in dataset_root.iterdir() if d.is_dir()]
classes = sorted(classes)
print("Found classes:", classes)
```

Found classes: ['cardboard', 'glass', 'metal', 'paper', 'plastic', 'trash']

## Step 3: Confirm expected classes & label consistency

```
In [ ]: expected = {"plastic", "metal", "glass", "cardboard", "paper", "trash"}
found = set(classes)
print("Expected classes:", expected)
print("Found classes:", found)

if found == expected:
    print("The classes match exactly the expected ones.")
else:
    print("Class mismatch.")
    print("Missing:", expected - found)
    print("Extra:", found - expected)
```

Expected classes: {'plastic', 'glass', 'paper', 'metal', 'cardboard', 'trash'}  
Found classes: {'plastic', 'glass', 'paper', 'metal', 'cardboard', 'trash'}  
The classes match exactly the expected ones.

## Step 4: Count images per class & class balance

```
In [ ]: # Supported Extensions
SUPPORTED_EXTS = (".jpg", ".jpeg", ".png")

class_counts = {}
for cls in classes:
    cls_dir = dataset_root / cls
    imgs = []
    for ext in SUPPORTED_EXTS:
        imgs.extend(list(cls_dir.glob(f"*{ext}")))
    # Also check for any unexpected extensions
    other = list(cls_dir.glob("*"))
    others = [p for p in other if p.suffix.lower() not in SUPPORTED_EXTS]
```

```

if others:
    print(f"Warning: found {len(others)} files in {cls} with unexpected extensions")
    class_counts[cls] = len(imgs)

print("Counts per class:")
for cls, cnt in class_counts.items():
    print(f" {cls}: {cnt}")
counts = np.array(list(class_counts.values()))
print("Total images:", counts.sum())
print("Min , Max , Mean:", counts.min(), ",", counts.max(), ",", counts.mean())

```

Counts per class:

```

cardboard: 2214
glass: 2500
metal: 2084
paper: 2315
plastic: 2288
trash: 2500

```

Total images: 13901

Min , Max , Mean: 2084 , 2500 , 2316.8333333333335

## Step 5: Duplicate Detection, Preview & Cleaned Copy (change path in code block)

```

In [ ]: import shutil
        from imagededup.methods import PHash

        # 1. Run imagededup PHash per class
        ph = PHash()
        dups_all = {}

        for cls in classes:
            cls_dir = dataset_root / cls
            print(f"Encoding class: {cls}")
            encodings = ph.encode_images(image_dir=str(cls_dir), recursive=False)
            dups = ph.find_duplicates(encoding_map=encodings, max_distance_threshold=3)
            dups_all[cls] = dups
            n_dup_keys = len([k for k, v in dups.items() if v])
            print(f" {n_dup_keys} keys have duplicates in {cls}")

        # 2. Collect duplicate pairs across all classes
        def collect_duplicate_pairs(dups_dict, cls):
            pairs = []
            for fname, dup_list in dups_dict.items():
                for dup in dup_list:
                    pairs.append((cls, fname, dup))
            return pairs

        all_pairs = []
        for cls, dups in dups_all.items():
            all_pairs.extend(collect_duplicate_pairs(dups, cls))

        print(f"\nTotal duplicate pairs found: {len(all_pairs)}")

```

```

# 3. Preview first 5 duplicate pairs
def preview_duplicate_pairs(pairs, n=5):
    for idx, (cls, f1, f2) in enumerate(pairs[:n]):
        path1 = dataset_root / cls / f1
        path2 = dataset_root / cls / f2

        fig, axes = plt.subplots(1, 2, figsize=(6, 3))
        try:
            img1 = Image.open(path1)
            img2 = Image.open(path2)

            axes[0].imshow(img1)
            axes[0].set_title(f"{f1}", fontsize=8)
            axes[0].axis("off")

            axes[1].imshow(img2)
            axes[1].set_title(f"{f2}", fontsize=8)
            axes[1].axis("off")

            plt.suptitle(f"Class: {cls} - Duplicate Pair {idx}")
            plt.show()
        except Exception as e:
            print(f"Error loading {f1}, {f2}:", e)

preview_duplicate_pairs(all_pairs, n=5)

# 4. Build cleaned dataset copy
cleaned_root = Path("data/Garbage_Dataset_Classification/images_cleaned")
cleaned_root.mkdir(parents=True, exist_ok=True)

# Mark duplicates to remove (always second file in pair)
to_remove = set([p[2] for p in all_pairs])
print("Total duplicate files to remove:", len(to_remove))

# Before counts
print("\nClass counts BEFORE cleaning:")
for cls in classes:
    total = sum(len(list((dataset_root/cls).glob(f"*{ext}"))) for ext in SUPPORTED_EXTENSIONS)
    print(f" {cls}: {total}")

# Copy files, skipping duplicates
for cls in classes:
    src_dir = dataset_root / cls
    dst_dir = cleaned_root / cls
    dst_dir.mkdir(parents=True, exist_ok=True)

    for ext in SUPPORTED_EXTENSIONS:
        for file in src_dir.glob(f"*{ext}"):
            if file.name not in to_remove:
                shutil.copy(file, dst_dir / file.name)

# After counts
print("\nClass counts AFTER cleaning:")
for cls in classes:
    total = sum(len(list((cleaned_root/cls).glob(f"*{ext}"))) for ext in SUPPORTED_EXTENSIONS)
    print(f" {cls}: {total}")

```

```
print("\n Cleaned dataset created at:", cleaned_root)
```

2025-10-12 15:51:25,446: INFO Start: Calculating hashes...

INFO:imagededup.methods.hashing:Start: Calculating hashes...

Encoding class: cardboard

100%|██████████| 2214/2214 [00:05<00:00, 393.12it/s]

2025-10-12 15:51:31,372: INFO End: Calculating hashes!

INFO:imagededup.methods.hashing:End: Calculating hashes!

/usr/local/lib/python3.12/dist-packages/imagededup/methods/hashing.py:317: RuntimeWarning: Parameter num\_enc\_workers has no effect since encodings are already provided  
warnings.warn('Parameter num\_enc\_workers has no effect since encodings are already provided', RuntimeWarning)

2025-10-12 15:51:31,374: INFO Start: Evaluating hamming distances for getting duplicates

INFO:imagededup.methods.hashing:Start: Evaluating hamming distances for getting duplicates

2025-10-12 15:51:31,376: INFO Start: Retrieving duplicates using Cython Brute force algorithm

INFO:imagededup.handlers.search.retrieval:Start: Retrieving duplicates using Cython Brute force algorithm

100%|██████████| 2214/2214 [00:04<00:00, 483.53it/s]

2025-10-12 15:51:36,045: INFO End: Retrieving duplicates using Cython Brute force algorithm

INFO:imagededup.handlers.search.retrieval:End: Retrieving duplicates using Cython Brute force algorithm

2025-10-12 15:51:36,049: INFO End: Evaluating hamming distances for getting duplicates

INFO:imagededup.methods.hashing:End: Evaluating hamming distances for getting duplicates

2025-10-12 15:51:36,105: INFO Start: Calculating hashes...

INFO:imagededup.methods.hashing:Start: Calculating hashes...

1071 keys have duplicates in cardboard

Encoding class: glass



```
100%|██████████| 2500/2500 [00:07<00:00, 328.41it/s]
2025-10-12 15:51:44,116: INFO End: Calculating hashes!
INFO:imagededup.methods.hashing:End: Calculating hashes!
2025-10-12 15:51:44,126: INFO Start: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:Start: Evaluating hamming distances for getting duplicates
2025-10-12 15:51:44,131: INFO Start: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:Start: Retrieving duplicates using Cython Brute force algorithm
100%|██████████| 2500/2500 [00:03<00:00, 755.09it/s]
2025-10-12 15:51:47,524: INFO End: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:End: Retrieving duplicates using Cython Brute force algorithm
2025-10-12 15:51:47,526: INFO End: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:End: Evaluating hamming distances for getting duplicates
2025-10-12 15:51:47,548: INFO Start: Calculating hashes...
INFO:imagededup.methods.hashing:Start: Calculating hashes...
```

```
591 keys have duplicates in glass
Encoding class: metal
```

```
100%|██████████| 2084/2084 [00:03<00:00, 568.17it/s]
2025-10-12 15:51:51,457: INFO End: Calculating hashes!
INFO:imagededup.methods.hashing:End: Calculating hashes!
2025-10-12 15:51:51,461: INFO Start: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:Start: Evaluating hamming distances for getting duplicates
2025-10-12 15:51:51,463: INFO Start: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:Start: Retrieving duplicates using Cython Brute force algorithm
100%|██████████| 2084/2084 [00:01<00:00, 1102.23it/s]
2025-10-12 15:51:53,420: INFO End: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:End: Retrieving duplicates using Cython Brute force algorithm
2025-10-12 15:51:53,421: INFO End: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:End: Evaluating hamming distances for getting duplicates
2025-10-12 15:51:53,445: INFO Start: Calculating hashes...
INFO:imagededup.methods.hashing:Start: Calculating hashes...
```

```
925 keys have duplicates in metal
Encoding class: paper
```

```
100%|██████████| 2315/2315 [00:02<00:00, 779.63it/s]
2025-10-12 15:51:56,565: INFO End: Calculating hashes!
INFO:imagededup.methods.hashing:End: Calculating hashes!
2025-10-12 15:51:56,569: INFO Start: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:Start: Evaluating hamming distances for getting duplicates
2025-10-12 15:51:56,570: INFO Start: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:Start: Retrieving duplicates using Cython Brute force algorithm
100%|██████████| 2315/2315 [00:01<00:00, 1265.60it/s]
2025-10-12 15:51:58,465: INFO End: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:End: Retrieving duplicates using Cython Brute force algorithm
2025-10-12 15:51:58,467: INFO End: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:End: Evaluating hamming distances for getting duplicates
2025-10-12 15:51:58,491: INFO Start: Calculating hashes...
INFO:imagededup.methods.hashing:Start: Calculating hashes...
```

1150 keys have duplicates in paper  
Encoding class: plastic

```
100%|██████████| 2288/2288 [00:02<00:00, 798.16it/s]
2025-10-12 15:52:01,557: INFO End: Calculating hashes!
INFO:imagededup.methods.hashing:End: Calculating hashes!
2025-10-12 15:52:01,560: INFO Start: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:Start: Evaluating hamming distances for getting duplicates
2025-10-12 15:52:01,561: INFO Start: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:Start: Retrieving duplicates using Cython Brute force algorithm
100%|██████████| 2288/2288 [00:03<00:00, 636.33it/s]
2025-10-12 15:52:05,227: INFO End: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:End: Retrieving duplicates using Cython Brute force algorithm
2025-10-12 15:52:05,228: INFO End: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:End: Evaluating hamming distances for getting duplicates
2025-10-12 15:52:05,275: INFO Start: Calculating hashes...
INFO:imagededup.methods.hashing:Start: Calculating hashes...
```

1309 keys have duplicates in plastic  
Encoding class: trash

```
100%|██████████| 2500/2500 [00:03<00:00, 769.52it/s]
2025-10-12 15:52:08,706: INFO End: Calculating hashes!
INFO:imagededup.methods.hashing:End: Calculating hashes!
2025-10-12 15:52:08,709: INFO Start: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:Start: Evaluating hamming distances for getting duplicates
2025-10-12 15:52:08,710: INFO Start: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:Start: Retrieving duplicates using Cython Brute force algorithm
100%|██████████| 2500/2500 [00:02<00:00, 1165.04it/s]
2025-10-12 15:52:10,908: INFO End: Retrieving duplicates using Cython Brute force algorithm
INFO:imagededup.handlers.search.retrieval:End: Retrieving duplicates using Cython Brute force algorithm
2025-10-12 15:52:10,909: INFO End: Evaluating hamming distances for getting duplicates
INFO:imagededup.methods.hashing:End: Evaluating hamming distances for getting duplicates
```

4 keys have duplicates in trash

Total duplicate pairs found: 5232

### Class: cardboard — Duplicate Pair 0

cardboard\_02078.jpg




cardboard\_00748.jpg



### Class: cardboard — Duplicate Pair 1

cardboard\_00791.jpg



cardboard\_02193.jpg



### Class: cardboard — Duplicate Pair 2

cardboard\_01980.jpg




cardboard\_00586.jpg



### Class: cardboard — Duplicate Pair 3

cardboard\_02226.jpg



cardboard\_01903.jpg



### Class: cardboard — Duplicate Pair 4

cardboard\_02840.jpg



cardboard\_01857.jpg



Total duplicate files to remove: 5050

Class counts BEFORE cleaning:

```
cardboard: 2214
glass: 2500
metal: 2084
paper: 2315
plastic: 2288
trash: 2500
```

Class counts AFTER cleaning:

```
cardboard: 1143
glass: 1909
metal: 1159
paper: 1165
plastic: 979
trash: 2496
```

Cleaned dataset created at: data/Garbage\_Dataset\_Classification/images\_cleaned

## Step 6: Confirm Image Sizes and Color Channels

```
In [ ]: shape_counter = Counter()
channel_counter = Counter()
bad_images = []

for cls in classes:
    cls_dir = cleaned_root / cls
    for ext in SUPPORTED_EXTS:
        for p in cls_dir.glob(f"*{ext}"):
            try:
                with Image.open(p) as img:
                    arr = np.array(img)
                    shape_counter[arr.shape] += 1
                    if arr.ndim == 3:
                        channel_counter[arr.shape[2]] += 1
                    else:
                        channel_counter[1] += 1
            except Exception as e:
                bad_images.append((p, str(e)))

print("Image shape distribution (H, W, [C]):")
for shp, cnt in shape_counter.items():
    print(f" {shp}: {cnt} images")

print("\nChannel counts:")
for c, cnt in channel_counter.items():
    print(f" {c} channels: {cnt} images")

print("\nNumber of images that failed to load:", len(bad_images))
if bad_images:
    print("Sample failed images:", bad_images[:5])
```

Image shape distribution (H, W, [C]):  
(256, 256, 3): 8851 images

Channel counts:  
3 channels: 8851 images

Number of images that failed to load: 0

## Step 7: Verify Labels Align with Images (using metadata.csv)

```
In [ ]: import pandas as pd

# 1. Verify every image in cleaned_root is inside the expected class folder
problems = []
for cls in classes:
    cls_dir = cleaned_root / cls
    for ext in SUPPORTED_EXTS:
        for p in cls_dir.glob(f"*{ext}"):
            if p.parent.name != cls:
                problems.append((p, p.parent.name, cls))

if problems:
    print(f"Found {len(problems)} images in the wrong folder:")
    print(problems[:10])
else:
    print("All images are in the correct class folders.")

# 2. Cross-check with metadata.csv located at the dataset's parent folder
metadata_path = dataset_root.parent / "metadata.csv"

if metadata_path.exists():
    meta = pd.read_csv(metadata_path)
    meta_map = dict(zip(meta.filename, meta.label))

    mismatches = []
    for cls in classes:
        cls_dir = cleaned_root / cls
        for ext in SUPPORTED_EXTS:
            for p in cls_dir.glob(f"*{ext}"):
                fname = p.name
                if fname in meta_map and meta_map[fname] != cls:
                    mismatches.append((fname, cls, meta_map[fname]))

    if mismatches:
        print(f"Found {len(mismatches)} mismatches with metadata.csv:")
        print(mismatches[:10])
    else:
        print("Folder labels match metadata.csv for all files checked.")
else:
    print("metadata.csv not found at:", metadata_path)
```

All images are in the correct class folders.  
Folder labels match metadata.csv for all files checked.

## Step 8: Stratified Split using StratifiedShuffleSplit (80/10/10)

```
In [ ]: from sklearn.model_selection import StratifiedShuffleSplit
        from collections import Counter

        # Rebuild data list from cleaned dataset
        data = []
        for cls in classes:
            cls_dir = cleaned_root / cls
            for ext in SUPPORTED_EXTS:
                for p in cls_dir.glob(f"*{ext}"):
                    data.append((p, cls))

        print("Total cleaned samples:", len(data))

        paths = [p for p, lbl in data]
        labels = [lbl for p, lbl in data]

        # Choose your split ratios
        test_ratio = 0.10
        val_ratio = 0.10
        train_ratio = 1.0 - (test_ratio + val_ratio)
        assert train_ratio > 0, "Make sure ratios sum to less than 1"

        # 1. Split off test set
        sss = StratifiedShuffleSplit(n_splits=1, test_size=test_ratio, random_state=42)
        for train_valid_idx, test_idx in sss.split(paths, labels):
            pass

        train_valid_paths = [paths[i] for i in train_valid_idx]
        train_valid_labels = [labels[i] for i in train_valid_idx]
        test_paths = [paths[i] for i in test_idx]
        test_labels = [labels[i] for i in test_idx]

        # 2. Split train_valid into train + validation
        rel_val = val_ratio / (train_ratio + val_ratio)
        sss2 = StratifiedShuffleSplit(n_splits=1, test_size=rel_val, random_state=42)
        for train_idx2, val_idx in sss2.split(train_valid_paths, train_valid_labels):
            pass

        train_paths = [train_valid_paths[i] for i in train_idx2]
        train_labels = [train_valid_labels[i] for i in train_idx2]
        val_paths = [train_valid_paths[i] for i in val_idx]
        val_labels = [train_valid_labels[i] for i in val_idx]

        # Build final splits
        train_set = list(zip(train_paths, train_labels))
        valid_set = list(zip(val_paths, val_labels))
        test_set = list(zip(test_paths, test_labels))
```

```

# Print sizes
print("Total:", len(data))
print("Train:", len(train_set), "Validation:", len(valid_set), "Test:", len(test_set))
print()

def print_dist(split, name):
    c = Counter(lbl for _, lbl in split)
    print(f"{name} class counts:")
    for cls in classes:
        print(f" {cls}: {c.get(cls, 0)}")
    print()

print_dist(train_set, "Train")
print_dist(valid_set, "Validation")
print_dist(test_set, "Test")

```

Total cleaned samples: 8851

Total: 8851

Train: 7079 Validation: 886 Test: 886

Train class counts:

```

cardboard: 915
glass: 1527
metal: 927
paper: 931
plastic: 783
trash: 1996

```

Validation class counts:

```

cardboard: 114
glass: 191
metal: 116
paper: 117
plastic: 98
trash: 250

```

Test class counts:

```

cardboard: 114
glass: 191
metal: 116
paper: 117
plastic: 98
trash: 250

```

## Step 9: Build PyTorch dataset & transforms (normalize + augment)

In [ ]: !pip install -q torch torchvision

```

import torch
from torch.utils.data import Dataset, DataLoader, WeightedRandomSampler
from torchvision import transforms, models

```



```

from PIL import Image
import numpy as np
import matplotlib.pyplot as plt
from collections import Counter
import os

# Map class names to integer labels (sorted to be consistent)
class_to_idx = {c: i for i, c in enumerate(classes)}
idx_to_class = {v: k for k, v in class_to_idx.items()}

# Transforms
# Train: light augmentations + normalize (ImageNet mean/std)
train_tfms = transforms.Compose([
    transforms.RandomHorizontalFlip(p=0.5),
    transforms.RandomRotation(degrees=10),
    transforms.RandomResizedCrop(size=224, scale=(0.9, 1.0)),
    transforms.ColorJitter(brightness=0.1, contrast=0.1),
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.485, 0.456, 0.406),
                          std=(0.229, 0.224, 0.225)),
])

# Val/Test: center crop + normalize (no augmentation)
eval_tfms = transforms.Compose([
    transforms.Resize(256),
    transforms.CenterCrop(224),
    transforms.ToTensor(),
    transforms.Normalize(mean=(0.485, 0.456, 0.406),
                          std=(0.229, 0.224, 0.225)),
])

class GarbageDataset(Dataset):
    def __init__(self, items, transform=None):
        """
        items: list of (Path, class_name)
        """
        self.items = items
        self.transform = transform

    def __len__(self):
        return len(self.items)

    def __getitem__(self, idx):
        path, cls_name = self.items[idx]
        label = class_to_idx[cls_name]
        img = Image.open(path).convert("RGB")
        if self.transform:
            img = self.transform(img)
        return img, label, str(path)

```

**Step 10: Address class imbalance after cleaning (class weights or weighted sampler)**

```
In [ ]: # Compute class counts from the TRAIN split
train_class_counts = Counter(lbl for _, lbl in train_set)
print("Train class counts:", train_class_counts)

# Option A: Class weights for CrossEntropyLoss: weight the loss function so false p
num_classes = len(classes)
counts = np.array([train_class_counts[c] for c in classes], dtype=np.float32)
class_weights = counts.sum() / (num_classes * counts) # inverse-frequency-ish
class_weights_tensor = torch.tensor(class_weights, dtype=torch.float32)
print("Class weights (A):", class_weights)

# Option B: WeightedRandomSampler: classes with fewer sample are more likely to be
label_to_idx = class_to_idx
per_class_weight = {cls: (counts.sum() / (num_classes * cnt))
                    for cls, cnt in train_class_counts.items()}
sample_weights = [per_class_weight[lbl] for _, lbl in train_set]
sampler = WeightedRandomSampler(weights=torch.DoubleTensor(sample_weights),
                                num_samples=len(sample_weights),
                                replacement=True)
```

Train class counts: Counter({'trash': 1996, 'glass': 1527, 'paper': 931, 'metal': 927, 'cardboard': 915, 'plastic': 783})  
Class weights (A): [1.2894354 0.7726479 1.2727436 1.2672753 1.5068114 0.5910988 5]

## Step 11: DataLoaders (with augmentation on train)

```
In [ ]: # Datasets
train_ds = GarbageDataset(train_set, transform=train_tfms)
val_ds = GarbageDataset(valid_set, transform=eval_tfms)
test_ds = GarbageDataset(test_set, transform=eval_tfms)

# DataLoaders: manage sampler, batch size, etc.
BATCH_SIZE = 32
num_workers = 2 if "COLAB_GPU" in os.environ or "COLAB_TPU_ADDR" in os.environ else

use_weighted_sampler = True # set False if we want to do option A

if use_weighted_sampler:
    train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE,
                              sampler=sampler, num_workers=num_workers, pin_memory=
else:
    train_loader = DataLoader(train_ds, batch_size=BATCH_SIZE,
                              shuffle=True, num_workers=num_workers, pin_memory=True)

val_loader = DataLoader(val_ds, batch_size=BATCH_SIZE,
                        shuffle=False, num_workers=num_workers, pin_memory=True)
test_loader = DataLoader(test_ds, batch_size=BATCH_SIZE,
                          shuffle=False, num_workers=num_workers, pin_memory=True)

print("Batches -> train:", len(train_loader), "val:", len(val_loader), "test:", len
```

Batches -> train: 222 val: 28 test: 28

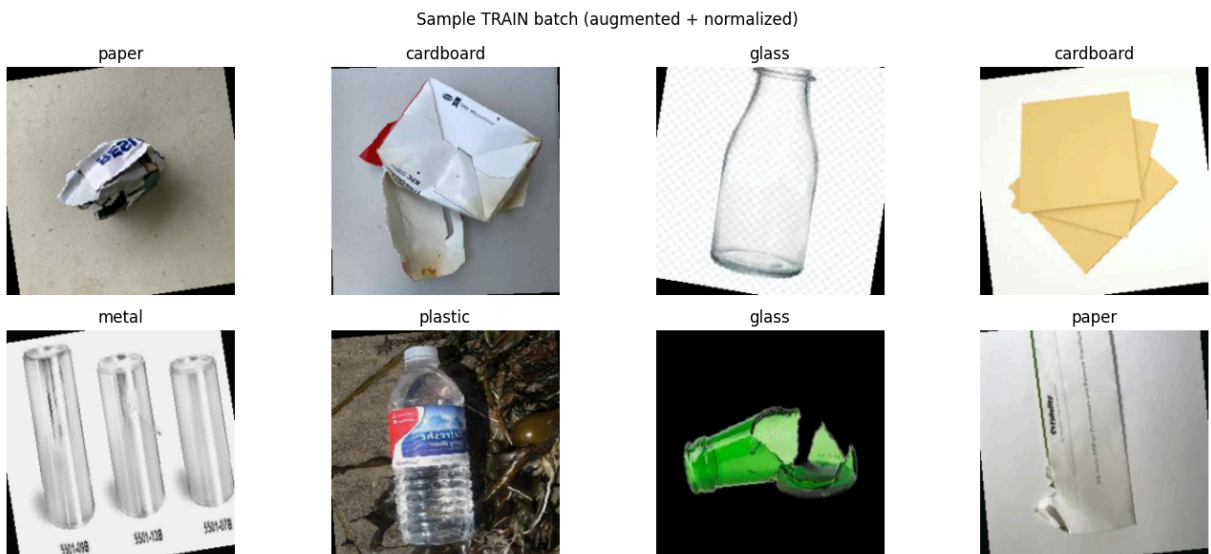
## Step 12: Visualize one preprocessed + augmented batch

```
In [ ]: # Helper to denormalize ImageNet-normalized tensors for display
IMAGENET_MEAN = np.array([0.485, 0.456, 0.406])
IMAGENET_STD = np.array([0.229, 0.224, 0.225])

def denormalize(img_tensor):
    img = img_tensor.detach().cpu().numpy().transpose(1,2,0)
    img = (img * IMAGENET_STD) + IMAGENET_MEAN
    img = np.clip(img, 0, 1)
    return img

# Get one batch
imgs, labels, paths = next(iter(train_loader))

# Plot first 8
n_show = min(8, imgs.size(0))
plt.figure(figsize=(14, 6))
for i in range(n_show):
    plt.subplot(2, 4, i+1)
    plt.imshow(denormalize(imgs[i]))
    plt.title(idx_to_class[int(labels[i])])
    plt.axis("off")
plt.suptitle("Sample TRAIN batch (augmented + normalized)")
plt.tight_layout()
plt.show()
```



## Step 13: Document split shapes (train/val/test per class)

```
In [ ]: def class_dist(items, title):
        c = Counter(lbl for _, lbl in items)
        print(title)
        for cls in classes:
            print(f" {cls:9s}: {c.get(cls,0)}")
        print(" TOTAL      :", sum(c.values()))
        print()

        class_dist(train_set, "TRAIN distribution")
        class_dist(valid_set, "VALID distribution")
        class_dist(test_set, "TEST distribution")
```

TRAIN distribution

```
cardboard: 915
glass      : 1527
metal      : 927
paper      : 931
plastic    : 783
trash      : 1996
TOTAL      : 7079
```

VALID distribution

```
cardboard: 114
glass      : 191
metal      : 116
paper      : 117
plastic    : 98
trash      : 250
TOTAL      : 886
```

TEST distribution

```
cardboard: 114
glass      : 191
metal      : 116
paper      : 117
plastic    : 98
trash      : 250
TOTAL      : 886
```

## Step 14: Pipeline readiness: push one batch through ResNet18

```
In [ ]: # Load a baseline CNN and run a single forward pass to confirm pipeline is ready
        device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
        print("Device:", device)

        # starting from scratch
        # later we can use pretrained weights for TL
        model = models.resnet18(weights=None)
        # Adapt the final layer to 6 classes
        model.fc = torch.nn.Linear(model.fc.in_features, len(classes))
        model = model.to(device)
```

```
# One forward pass
model.eval()
with torch.no_grad():
    xb, yb, _ = next(iter(train_loader))
    xb = xb.to(device)
    yb = yb.to(device)
    logits = model(xb)
    print("Forward OK -> logits shape:", logits.shape)
```

Device: cuda

Forward OK -> logits shape: torch.Size([32, 6])

## Step 15: MobileNetV2 Architecture

```
In [ ]: from torch import nn, optim
        from torch.optim.lr_scheduler import CosineAnnealingLR
        from tqdm.notebook import tqdm

        # Depthwise Separable Convolution (used inside rInverted Residual)
        class Conv_BN_ReLU(nn.Sequential):
            def __init__(self, in_channels, out_channels, kernel_size=3, stride=1, groups=1):

                padding = (kernel_size - 1) // 2 # to preserve original size if stride = 1

                super().__init__(
                    # If group = 1, it's standard convolution (1 filter for entire depth)
                    # if groups = in_channels it's depthwise convolution (1 filter per layer)
                    nn.Conv2d(in_channels, out_channels, kernel_size, stride, padding, groups=groups),
                    nn.BatchNorm2d(out_channels),
                    nn.ReLU6(inplace=True) # ReLU that clips value to be between [0, 6]
                )

        # Inverted Residual Block
        class Inverted_Residual(nn.Module):
            def __init__(self, in_channels, out_channels, stride, expand_ratio):
                super().__init__()
                # expand ratio control how much we expand channels before depthwise conv.
                '''more channels/hidden_dim means more depth/layers or feature maps produce'''
                hidden_dim = in_channels * expand_ratio

                '''a boolean that determine if residual short cut is applied at the end'''
                self.use_res_connect = (stride == 1 and in_channels == out_channels)

                layers = []
                if expand_ratio != 1:
                    # carry out the expansion from in_channels to hidden_dim size if expand
                    layers.append(Conv_BN_ReLU(in_channels, hidden_dim, kernel_size=1))
                    '''Narrow -> Wide (recall Invert Residual Block is: Narrow -> Wide -> N

                # applies Depthwise conv, 3x3
                '''since groups=in_channels it means 1 filter per layer = Depthwise conv'''
                layers.append(Conv_BN_ReLU(hidden_dim, hidden_dim, stride=stride, groups=hi
```

```

# applies 1x1 conv: communication between depths
'''compress channels back to low-dimensional space (out_channels)'''
'''Wide -> Narrow'''
layers.append(nn.Conv2d(hidden_dim, out_channels, kernel_size=1, stride=1,
layers.append(nn.BatchNorm2d(out_channels))
# no ReLU because of linear bottleneck (prevents losing information in comp

'''the *layer is same as passing individual element of layer array, it's ju
self.conv = nn.Sequential(*layers)
# warps the layers into a sequential block

def forward(self, x):
    if self.use_res_connect:
        # if dimension of in and out channel match at the start
        # adds input to the output here at the end (residual short cut connecti
        return x + self.conv(x)
    else:
        return self.conv(x)

class MobileNetV2(nn.Module):
    # width_mult to scale channels up/down
    def __init__(self, num_classes=6, width_mult=1.0, dropout_rate=0.2):
        super().__init__()

        inverted_residual_setting = [
            # t (expand ratio), c (channels), n (repeats), s (stride)
            [1, 16, 1, 1],
            [6, 24, 2, 2],
            [6, 32, 3, 2],
            [6, 64, 4, 2],
            [6, 96, 3, 1],
            [6, 160, 3, 2],
            [6, 320, 1, 1],
        ]
        ...

        t: how much to widen channels inside the block (more channel=more layer
        c: number of channels after each projection (output channels)
        n: number of times to repeat the blocks
        s: stride of first block, rest block all have stride = 1 (meaning down
        ...

        input_channel = int(32 * width_mult)
        last_channel = int(1280*width_mult) if width_mult > 1.0 else 1280
        # first conv layer outputs 32 channel scaled by width_mult
        # final conv layer outputs 1280 channel scaled

        '''add first standard 3x3 conv from RGB of 3 channel to 32 channel scaled''
        features = [Conv_BN_ReLU(in_channels=3, out_channels=input_channel, stride=

        '''loop through each stage Continuously adding layers: Expands -> depthwise
        apply stride in first block = downsampling, then rest block keep stride=
        for t, c, n, s in inverted_residual_setting:
            output_channel = int(c * width_mult) # Expand
            for i in range(n):
                stride = s if i == 0 else 1 # stride of 1 if not first block
                features.append(Inverted_Residual(input_channel, output_channel, st

```

```

        input_channel = output_channel # set the corresponding input channel

    # add final 1x1 convolution
    features.append(Conv_BN_ReLU(input_channel, last_channel, kernel_size=1))

    # wrap all layer into a sequential block
    self.features = nn.Sequential(*features)

    # define final linear classification layer added in forward function
    self.classifier = nn.Sequential(
        nn.Dropout(dropout_rate),
        nn.Linear(last_channel, num_classes),
    )

    '''initialize weight before forward pass'''
    self._initialize_weights()

    '''define forward pass:
    1 extract features
    2 apply global average pooling
    3 fully connected layer for classification'''
    def forward(self, x):
        x = self.features(x) # continuous feature extraction
        x = x.mean([2, 3]) # Add global average pooling layer
        x = self.classifier(x) # Linear layer for final classification
        return x

    # the "_" before function indicates it's an function for internal use
    def _initialize_weights(self):
        '''loop through all layers'''
        for m in self.modules():
            '''check if current layer is either 2D conv. or 2D batch norm. or a full
            if isinstance(m, nn.Conv2d):
                # value are normally distributed mean of 0 with variance scaled based on
                nn.init.kaiming_normal_(m.weight, mode="fan_out")
                if m.bias is not None: # initialize bias if exist in current layer
                    nn.init.zeros_(m.bias)
            elif isinstance(m, nn.BatchNorm2d):
                nn.init.ones_(m.weight)
                nn.init.zeros_(m.bias)
            elif isinstance(m, nn.Linear):
                # weight initialize with mean 0 and deviation 0.01
                nn.init.normal_(m.weight, 0, 0.01)
                nn.init.zeros_(m.bias)

```

## Step 16: Training and Evaluation Functions

```

In [ ]: from sklearn.metrics import precision_recall_fscore_support, confusion_matrix, clas
import seaborn as sns

def model_train(model, train_loader, validation_loader, optimizer, criterion, sched
    # Tracking Lists
    epoch_list = []
    train_losses, val_losses = [], []

```

```

train_accuracies, val_accuracies = [], []
train_precisions, val_precisions = [], []
train_recalls, val_recalls = [], []
train_f1s, val_f1s = [], []

best_acc = 0.0
best_f1 = 0.0

for epoch in range(epochs):
    # ===== TRAINING PHASE =====
    model.train()
    running_loss = 0.0
    all_preds, all_labels = [], []

    for images, labels, paths in tqdm(train_loader, desc=f"Epoch {epoch+1}/{epochs}"):
        images, labels = images.to(device), labels.to(device)

        optimizer.zero_grad()
        y_pred = model(images)
        loss = criterion(y_pred, labels)
        loss.backward()
        optimizer.step()

        running_loss += loss.item() * images.size(0)
        _, predicted = torch.max(y_pred, 1)

        all_preds.extend(predicted.cpu().numpy())
        all_labels.extend(labels.cpu().numpy())

    # Calculate training metrics
    train_loss = running_loss / len(train_loader.dataset)
    train_acc = 100 * np.mean(np.array(all_preds) == np.array(all_labels))

    # Calculate precision, recall, F1
    precision, recall, f1, _ = precision_recall_fscore_support(
        all_labels, all_preds, average='macro', zero_division=0
    )

    # ===== VALIDATION PHASE =====
    val_metrics = model_evaluation(model, validation_loader, criterion, device)

    # Print epoch summary
    print(f"Epoch [{epoch+1}/{epochs}]")
    print(f"  Train -> Loss: {train_loss:.4f} | Acc: {train_acc:.2f}% | "
          f"Precision: {precision:.4f} | Recall: {recall:.4f} | F1: {f1:.4f}")
    print(f"  Val   -> Loss: {val_metrics['loss']:.4f} | Acc: {val_metrics['acc']:.2f}% | "
          f"Precision: {val_metrics['precision']:.4f} | Recall: {val_metrics['recall']:.4f} | "
          f"F1: {val_metrics['f1']:.4f}")

    # Store metrics
    epoch_list.append(epoch + 1)
    train_losses.append(train_loss)
    train_accuracies.append(train_acc)
    train_precisions.append(precision)
    train_recalls.append(recall)
    train_f1s.append(f1)

```



```

val_losses.append(val_metrics['loss'])
val_accuracies.append(val_metrics['accuracy'])
val_precisions.append(val_metrics['precision'])
val_recalls.append(val_metrics['recall'])
val_f1s.append(val_metrics['f1'])

scheduler.step()

# Save best model based on validation F1 score
if val_metrics['f1'] > best_f1:
    best_f1 = val_metrics['f1']
    best_acc = val_metrics['accuracy']
    torch.save({
        'epoch': epoch,
        'model_state_dict': model.state_dict(),
        'optimizer_state_dict': optimizer.state_dict(),
        'val_f1': best_f1,
        'val_acc': best_acc,
    }, 'best_model.pth')
    print(f" ✓ New best model saved! (F1: {best_f1:.4f}, Acc: {best_acc:.2f})")

return {
    'epochs': epoch_list,
    'train_losses': train_losses,
    'val_losses': val_losses,
    'train_accuracies': train_accuracies,
    'val_accuracies': val_accuracies,
    'train_precisions': train_precisions,
    'val_precisions': val_precisions,
    'train_recalls': train_recalls,
    'val_recalls': val_recalls,
    'train_f1s': train_f1s,
    'val_f1s': val_f1s,
    'best_f1': best_f1,
    'best_acc': best_acc
}

```

```

def model_evaluation(model, data_loader, criterion, device):
    model.eval()
    running_loss = 0.0
    all_preds, all_labels = [], []

    with torch.no_grad():
        for images, labels, paths in data_loader:
            images, labels = images.to(device), labels.to(device)
            y_pred = model(images)
            loss = criterion(y_pred, labels)

            running_loss += loss.item() * images.size(0)
            _, predicted = torch.max(y_pred, dim=1)

            all_preds.extend(predicted.cpu().numpy())
            all_labels.extend(labels.cpu().numpy())

```

```

# Calculate overall metrics
avg_loss = running_loss / len(data_loader.dataset)
accuracy = 100 * np.mean(np.array(all_preds) == np.array(all_labels))

# Calculate precision, recall, F1 (macro-averaged)
precision, recall, f1, _ = precision_recall_fscore_support(
    all_labels, all_preds, average='macro', zero_division=0
)

# Calculate per-class metrics
per_class_precision, per_class_recall, per_class_f1, support = precision_recall_fscore_support(
    all_labels, all_preds, average=None, zero_division=0
)

return {
    'loss': avg_loss,
    'accuracy': accuracy,
    'precision': precision,
    'recall': recall,
    'f1': f1,
    'per_class_precision': per_class_precision,
    'per_class_recall': per_class_recall,
    'per_class_f1': per_class_f1,
    'support': support,
    'all_preds': all_preds,
    'all_labels': all_labels
}

```

## Step 17: Hyperparameters & Training Configuration

```

In [ ]: # hyper parameter
learning_rate = 0.001
weight_decay = 1e-4
num_epochs = 10
dropout_rate = 0.2

if __name__ == "__main__":
    # Optional warning suppression
    os.environ['PYTHONWARNINGS'] = 'ignore:semaphore_tracker:UserWarning'
    # This ensures proper multiprocessing behavior
    # If using multiprocessing DataLoader
    import sys, torch.multiprocessing as mp
    if sys.platform.startswith("linux"):
        mp.set_start_method('fork', force=True) # works and fixed your case
    else:
        mp.set_start_method('spawn', force=True) # safer on macOS/Windows

    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    print(f"Using device: {device}")

```

```

model = MobileNetV2(num_classes=6, width_mult=1.0, dropout_rate=dropout_rate).to(device)
print(f"Model parameters: {sum(p.numel() for p in model.parameters()):,}")

# Optimizer, loss, scheduler
optimizer = optim.Adam(model.parameters(), lr=learning_rate, weight_decay=weight_decay)
criterion = nn.CrossEntropyLoss()
scheduler = torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=num_epochs)

# Train model
print(f"\nStarting training for {num_epochs} epochs...")
results = model_train(model, train_loader, val_loader, optimizer, criterion, scheduler)

print(f"\n{'='*60}")
print(f"Training Complete!")
print(f"Best Validation F1: {results['best_f1']:.4f}")
print(f"Best Validation Accuracy: {results['best_acc']:.2f}%")
print(f"{'='*60}")

```

Using device: cuda

Model parameters: 2,231,558

Starting training for 10 epochs...

Epoch 1/10: 0%| | 0/222 [00:00<?, ?it/s]

Epoch [1/10]

Train -> Loss: 1.5292 | Acc: 39.62% | Precision: 0.3919 | Recall: 0.3974 | F1: 0.3926

Val -> Loss: 1.4790 | Acc: 43.68% | Precision: 0.4978 | Recall: 0.4379 | F1: 0.4055

✓ New best model saved! (F1: 0.4055, Acc: 43.68%)

Epoch 2/10: 0%| | 0/222 [00:00<?, ?it/s]

Epoch [2/10]

Train -> Loss: 1.3327 | Acc: 48.26% | Precision: 0.4818 | Recall: 0.4821 | F1: 0.4805

Val -> Loss: 1.4054 | Acc: 47.29% | Precision: 0.4949 | Recall: 0.4703 | F1: 0.4233

✓ New best model saved! (F1: 0.4233, Acc: 47.29%)

Epoch 3/10: 0%| | 0/222 [00:00<?, ?it/s]

Epoch [3/10]

Train -> Loss: 1.2190 | Acc: 54.84% | Precision: 0.5484 | Recall: 0.5480 | F1: 0.5464

Val -> Loss: 1.2634 | Acc: 51.24% | Precision: 0.5480 | Recall: 0.5414 | F1: 0.5029

✓ New best model saved! (F1: 0.5029, Acc: 51.24%)

Epoch 4/10: 0%| | 0/222 [00:00<?, ?it/s]

Epoch [4/10]

Train -> Loss: 1.0852 | Acc: 60.57% | Precision: 0.6061 | Recall: 0.6056 | F1: 0.6052

Val -> Loss: 1.1457 | Acc: 56.55% | Precision: 0.6108 | Recall: 0.5418 | F1: 0.5424

✓ New best model saved! (F1: 0.5424, Acc: 56.55%)

Epoch 5/10: 0%| | 0/222 [00:00<?, ?it/s]

```

Epoch [5/10]
  Train -> Loss: 1.0088 | Acc: 63.40% | Precision: 0.6345 | Recall: 0.6336 | F1: 0.6330
  Val   -> Loss: 1.2023 | Acc: 58.80% | Precision: 0.6339 | Recall: 0.5864 | F1: 0.5620
    ✓ New best model saved! (F1: 0.5620, Acc: 58.80%)
Epoch 6/10:  0%|                | 0/222 [00:00<?, ?it/s]
Epoch [6/10]
  Train -> Loss: 0.9245 | Acc: 65.73% | Precision: 0.6558 | Recall: 0.6558 | F1: 0.6551
  Val   -> Loss: 0.9403 | Acc: 65.46% | Precision: 0.6514 | Recall: 0.6523 | F1: 0.6454
    ✓ New best model saved! (F1: 0.6454, Acc: 65.46%)
Epoch 7/10:  0%|                | 0/222 [00:00<?, ?it/s]
Epoch [7/10]
  Train -> Loss: 0.8349 | Acc: 70.14% | Precision: 0.7009 | Recall: 0.7011 | F1: 0.7008
  Val   -> Loss: 0.9722 | Acc: 64.33% | Precision: 0.6394 | Recall: 0.6293 | F1: 0.6230
Epoch 8/10:  0%|                | 0/222 [00:00<?, ?it/s]
Epoch [8/10]
  Train -> Loss: 0.7843 | Acc: 70.94% | Precision: 0.7091 | Recall: 0.7090 | F1: 0.7088
  Val   -> Loss: 0.8161 | Acc: 70.43% | Precision: 0.6909 | Recall: 0.7147 | F1: 0.6959
    ✓ New best model saved! (F1: 0.6959, Acc: 70.43%)
Epoch 9/10:  0%|                | 0/222 [00:00<?, ?it/s]
Epoch [9/10]
  Train -> Loss: 0.7211 | Acc: 73.64% | Precision: 0.7363 | Recall: 0.7371 | F1: 0.7363
  Val   -> Loss: 0.7475 | Acc: 73.14% | Precision: 0.7143 | Recall: 0.7316 | F1: 0.7208
    ✓ New best model saved! (F1: 0.7208, Acc: 73.14%)
Epoch 10/10: 0%|                | 0/222 [00:00<?, ?it/s]
Epoch [10/10]
  Train -> Loss: 0.6817 | Acc: 75.22% | Precision: 0.7522 | Recall: 0.7516 | F1: 0.7514
  Val   -> Loss: 0.7210 | Acc: 74.15% | Precision: 0.7227 | Recall: 0.7430 | F1: 0.7305
    ✓ New best model saved! (F1: 0.7305, Acc: 74.15%)

=====
Training Complete!
Best Validation F1: 0.7305
Best Validation Accuracy: 74.15%
=====

```

## Step 18: Visualize Metrics

```

In [ ]: # Plot training history
fig, axes = plt.subplots(2, 2, figsize=(15, 10))

# Plot 1: Loss over epochs
axes[0, 0].plot(results['epochs'], results['train_losses'], label='Train Loss', mar

```

```

axes[0, 0].plot(results['epochs'], results['val_losses'], label='Val Loss', marker=
axes[0, 0].set_xlabel('Epoch', fontsize=12)
axes[0, 0].set_ylabel('Loss', fontsize=12)
axes[0, 0].set_title('Training and Validation Loss', fontsize=14, fontweight='bold'
axes[0, 0].legend(fontsize=11)
axes[0, 0].grid(True, alpha=0.3)

# Plot 2: Accuracy over epochs
axes[0, 1].plot(results['epochs'], results['train_accuracies'], label='Train Accura
axes[0, 1].plot(results['epochs'], results['val_accuracies'], label='Val Accuracy',
axes[0, 1].set_xlabel('Epoch', fontsize=12)
axes[0, 1].set_ylabel('Accuracy (%)', fontsize=12)
axes[0, 1].set_title('Training and Validation Accuracy', fontsize=14, fontweight='b
axes[0, 1].legend(fontsize=11)
axes[0, 1].grid(True, alpha=0.3)

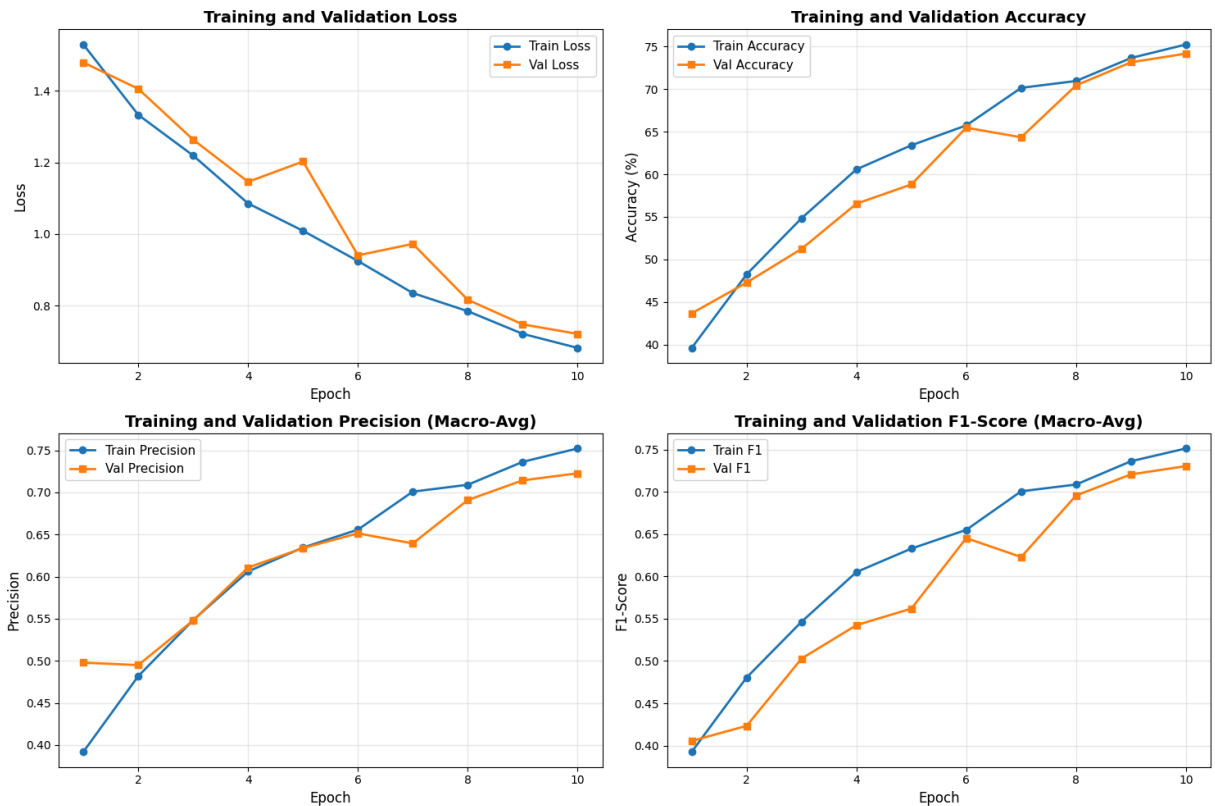
# Plot 3: Precision over epochs
axes[1, 0].plot(results['epochs'], results['train_precisions'], label='Train Preci
axes[1, 0].plot(results['epochs'], results['val_precisions'], label='Val Precision'
axes[1, 0].set_xlabel('Epoch', fontsize=12)
axes[1, 0].set_ylabel('Precision', fontsize=12)
axes[1, 0].set_title('Training and Validation Precision (Macro-Avg)', fontsize=14,
axes[1, 0].legend(fontsize=11)
axes[1, 0].grid(True, alpha=0.3)

# Plot 4: F1-Score over epochs
axes[1, 1].plot(results['epochs'], results['train_f1s'], label='Train F1', marker=
axes[1, 1].plot(results['epochs'], results['val_f1s'], label='Val F1', marker='s',
axes[1, 1].set_xlabel('Epoch', fontsize=12)
axes[1, 1].set_ylabel('F1-Score', fontsize=12)
axes[1, 1].set_title('Training and Validation F1-Score (Macro-Avg)', fontsize=14, f
axes[1, 1].legend(fontsize=11)
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()

# Print final results summary
print(f"\n{'='*60}")
print("FINAL TRAINING SUMMARY")
print(f"{'='*60}")
print(f"Best Validation Accuracy: {results['best_acc']:.2f}%")
print(f"Best Validation F1-Score: {results['best_f1']:.4f}")
print(f"\nFinal Epoch Metrics:")
print(f"  Train -> Loss: {results['train_losses'][-1]:.4f} | Acc: {results['train_a
print(f"  Val   -> Loss: {results['val_losses'][-1]:.4f} | Acc: {results['val_accu
print(f"{'='*60}")

```



```
=====
FINAL TRAINING SUMMARY
=====
Best Validation Accuracy: 74.15%
Best Validation F1-Score: 0.7305

Final Epoch Metrics:
  Train -> Loss: 0.6817 | Acc: 75.22% | F1: 0.7514
  Val   -> Loss: 0.7210 | Acc: 74.15% | F1: 0.7305
=====
```

## Step 19: Confusion Matrix Visualization

```
In [ ]: # Evaluate final model on validation set to get predictions for confusion matrix
final_val_metrics = model_evaluation(model, val_loader, criterion, device)

# Create confusion matrix
cm = confusion_matrix(final_val_metrics['all_labels'], final_val_metrics['all_preds'])

# Plot confusion matrix
fig, axes = plt.subplots(1, 2, figsize=(18, 7))

# Raw counts confusion matrix
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=classes, yticklabels=classes,
            ax=axes[0], cbar_kws={'label': 'Count'})
axes[0].set_xlabel('Predicted Label', fontsize=12, fontweight='bold')
axes[0].set_ylabel('True Label', fontsize=12, fontweight='bold')
axes[0].set_title('Confusion Matrix (Counts)', fontsize=14, fontweight='bold')
```

```


# Normalized confusion matrix (percentages)
cm_normalized = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]
sns.heatmap(cm_normalized, annot=True, fmt='.2%', cmap='Blues',
            xticklabels=classes, yticklabels=classes,
            ax=axes[1], cbar_kws={'label': 'Percentage'})
axes[1].set_xlabel('Predicted Label', fontsize=12, fontweight='bold')
axes[1].set_ylabel('True Label', fontsize=12, fontweight='bold')
axes[1].set_title('Confusion Matrix (Normalized by True Class)', fontsize=14, fontw

plt.tight_layout()
plt.show()

# Print classification report
print("\n" + "="*70)
print("CLASSIFICATION REPORT (Validation Set)")
print("="*70)
print(classification_report(final_val_metrics['all_labels'],
                            final_val_metrics['all_preds'],
                            target_names=classes,
                            digits=4))

# Print per-class metrics in a table format
print("\n" + "="*70)
print("PER-CLASS METRICS SUMMARY")
print("="*70)
print(f"{'Class':<12} {'Precision':<12} {'Recall':<12} {'F1-Score':<12} {'Support':<12}")
print("-"*70)
for i, cls in enumerate(classes):
    print(f"{'cls':<12} {final_val_metrics['per_class_precision'][i]:<12.4f} "
          f"{final_val_metrics['per_class_recall'][i]:<12.4f} "
          f"{final_val_metrics['per_class_f1'][i]:<12.4f} "
          f"{int(final_val_metrics['support'][i]):<10}")
print("-"*70)
print(f"{'Macro Avg':<12} {final_val_metrics['precision']:<12.4f} "
      f"{final_val_metrics['recall']:<12.4f} "
      f"{final_val_metrics['f1']:<12.4f} "
      f"{int(sum(final_val_metrics['support'])):<10}")
print("="*70)

```



=====				
CLASSIFICATION REPORT (Validation Set)				
=====				
	precision	recall	f1-score	support
-----				
cardboard	0.7619	0.8421	0.8000	114
glass	0.7865	0.7330	0.7588	191
metal	0.6850	0.7500	0.7160	116
paper	0.6815	0.7863	0.7302	117
plastic	0.5755	0.6224	0.5980	98
trash	0.8458	0.7240	0.7802	250
-----				
accuracy			0.7415	886
macro avg	0.7227	0.7430	0.7305	886
weighted avg	0.7496	0.7415	0.7430	886

=====				
PER-CLASS METRICS SUMMARY				
=====				
Class	Precision	Recall	F1-Score	Support
-----				
cardboard	0.7619	0.8421	0.8000	114
glass	0.7865	0.7330	0.7588	191
metal	0.6850	0.7500	0.7160	116
paper	0.6815	0.7863	0.7302	117
plastic	0.5755	0.6224	0.5980	98
trash	0.8458	0.7240	0.7802	250
-----				
Macro Avg	0.7227	0.7430	0.7305	886
=====				